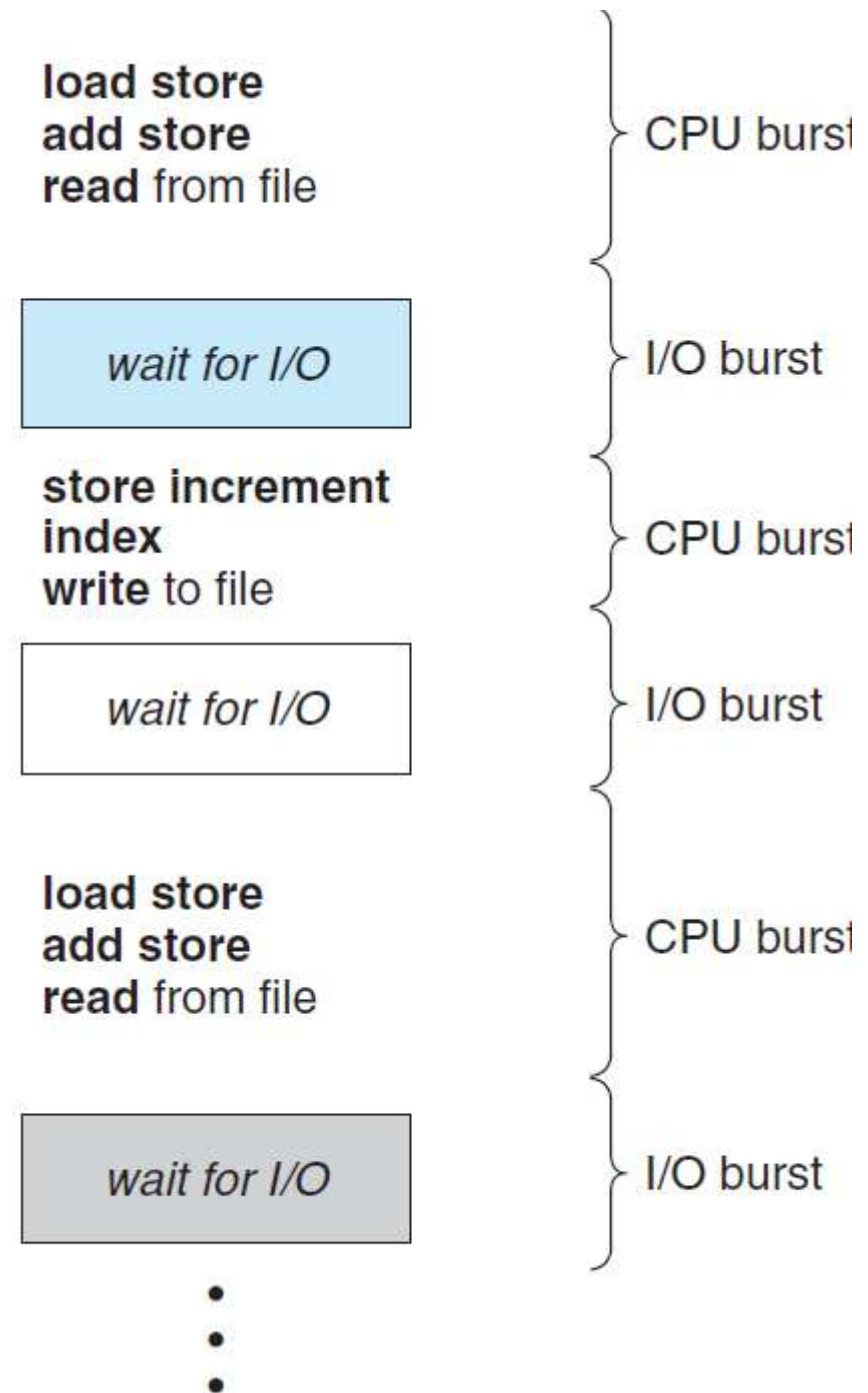


CPU Scheduling


Course Instructor: Nausheen Shoaib

Basic Concepts

- ▶ Maximum CPU utilization obtained with multiprogramming
- ▶ CPU-I/O Burst Cycle – Process execution consists of a **cycle** of CPU execution and I/O wait
- ▶ **CPU burst** followed by **I/O burst**
- ▶ CPU burst distribution is of main concern



Histogram of CPU-burst Times



Preemptive Vs. Non Preemptive Scheduling

1. The basic difference between preemptive and non-preemptive scheduling is that in preemptive scheduling the CPU is allocated to the processes for the **limited** time. While in Non-preemptive scheduling, the CPU is allocated to the process till it **terminates** or switches to **waiting state**.
2. The executing process in preemptive scheduling is **interrupted** in the middle of execution whereas, the executing process in non-preemptive scheduling is **not interrupted** in the middle of execution.
3. Preemptive Scheduling has the **overhead** of switching the process from ready state to running state, vise-verse, and maintaining the ready queue. On the other hands, non-preemptive scheduling has **no overhead** of switching the process from running state to ready state.
4. In preemptive scheduling, if a process with high priority frequently arrives in the ready queue then the process with low priority have to wait for a long, and it may have to starve. On the other hands, in the non-preemptive scheduling, if CPU is allocated to the process with larger burst time then the processes with small burst time may have to starve.
5. Preemptive scheduling is quite **flexible** because the critical processes are allowed to access CPU as they arrive into the ready queue, no matter what process is executing currently. Non-preemptive scheduling is **rigid** as even if a critical process enters the ready queue the process running CPU is not disturbed.


Dispatcher

- ▶ Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - switching context
 - switching to user mode
 - jumping to the proper location in the user program to restart that program
- ▶ **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running

Scheduling Criteria

- ▶ **CPU utilization** – keep the CPU as busy as possible
- ▶ **Throughput** – # of processes that complete their execution per time unit
- ▶ **Turnaround time** – amount of time to execute a particular process
- ▶ Turnaround time (TAT)=Completion time – Arrival time
- ▶ **Waiting time** – amount of time a process has been waiting in the ready queue
- ▶ **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

Scheduling Algorithm Optimization Criteria

- ▶ Max CPU utilization
 - ▶ Max throughput
 - ▶ Min turnaround time
 - ▶ Min waiting time
 - ▶ Min response time
- 

First-Come, First-Served (FCFS) Scheduling


Process Burst Time

P_1 24

P_2 3

P_3 3

- ▶ Suppose that the processes arrive in the order: P_1 , P_2 , P_3



Example FCFS


- ▶ Covered in class

FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

P_2, P_3, P_1

- ▶ The Gantt chart for the schedule is:



- ▶ Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- ▶ Average waiting time: $(6 + 0 + 3)/3 = 3$
- ▶ Much better than previous case
- ▶ **Convoy effect** – short process behind long process
 - Consider one CPU-bound and many I/O-bound processes


Shortest–Job–First (SJF) Scheduling

- ▶ Associate with each process the length of its next CPU burst
 - Use these lengths to schedule the process with the shortest time
- ▶ SJF is optimal – gives minimum average waiting time for a given set of processes
 - The difficulty is knowing the length of the next CPU request
 - Could ask the user

Example of SJF

<u>Process</u>	<u>Burst Time</u>
P_1	6
P_2	8
P_3	7
P_4	3

- ▶ SJF scheduling chart



- ▶ Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$

Example of SJF

- ▶ Covered in class

Determining Length of Next CPU Burst

- ▶ Can only estimate the length – should be similar to the previous one
 - Then pick process with shortest predicted next CPU burst
- ▶ Can be done by using the length of previous CPU bursts, using exponential averaging

$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n.$$

- ▶ t_n = most recent info
- ▶ T_n = Past history
- ▶ If $\alpha=0$ then recent history has no effect
- ▶ If $\alpha=1$, then =recent CPU burst matters
- ▶ Preemptive version called **shortest-remaining-time-first**

Shortest-remaining-time-first

- ▶ Now we add the concepts of varying arrival times and preemption to the analysis:
- ▶ the process with the smallest amount of time remaining until completion is selected to execute.

Example of SRTF

- ▶ Covered in class


Priority Scheduling

- ▶ A priority number (integer) is associated with each process
- ▶ The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - Preemptive
 - Non-preemptive
- ▶ Problem \equiv **Starvation** – low priority processes may never execute
- ▶ Solution \equiv **Aging** – as time progresses increase the priority of the process

Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

► Priority scheduling Gantt Chart



► Average waiting time = 8.2 msec

Example Priority

- ▶ Covered in class


Round Robin (RR)

- ▶ Each process gets a small unit of CPU time (**time quantum q**), usually 10–100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- ▶ If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- ▶ Timer interrupts every quantum to schedule next process
- ▶ Performance
 - q large \Rightarrow FIFO
 - q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high

Example of RR with Time Quantum = 4

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- ▶ The Gantt chart is:



- ▶ **Completion time** = $P1=30$; $P2=7$; $P3=10$
- ▶ **TAT** = $P1=(30-0)$; $P2=(7-4)$; $P3=(10-7)$
- ▶ **Avg. TAT**=

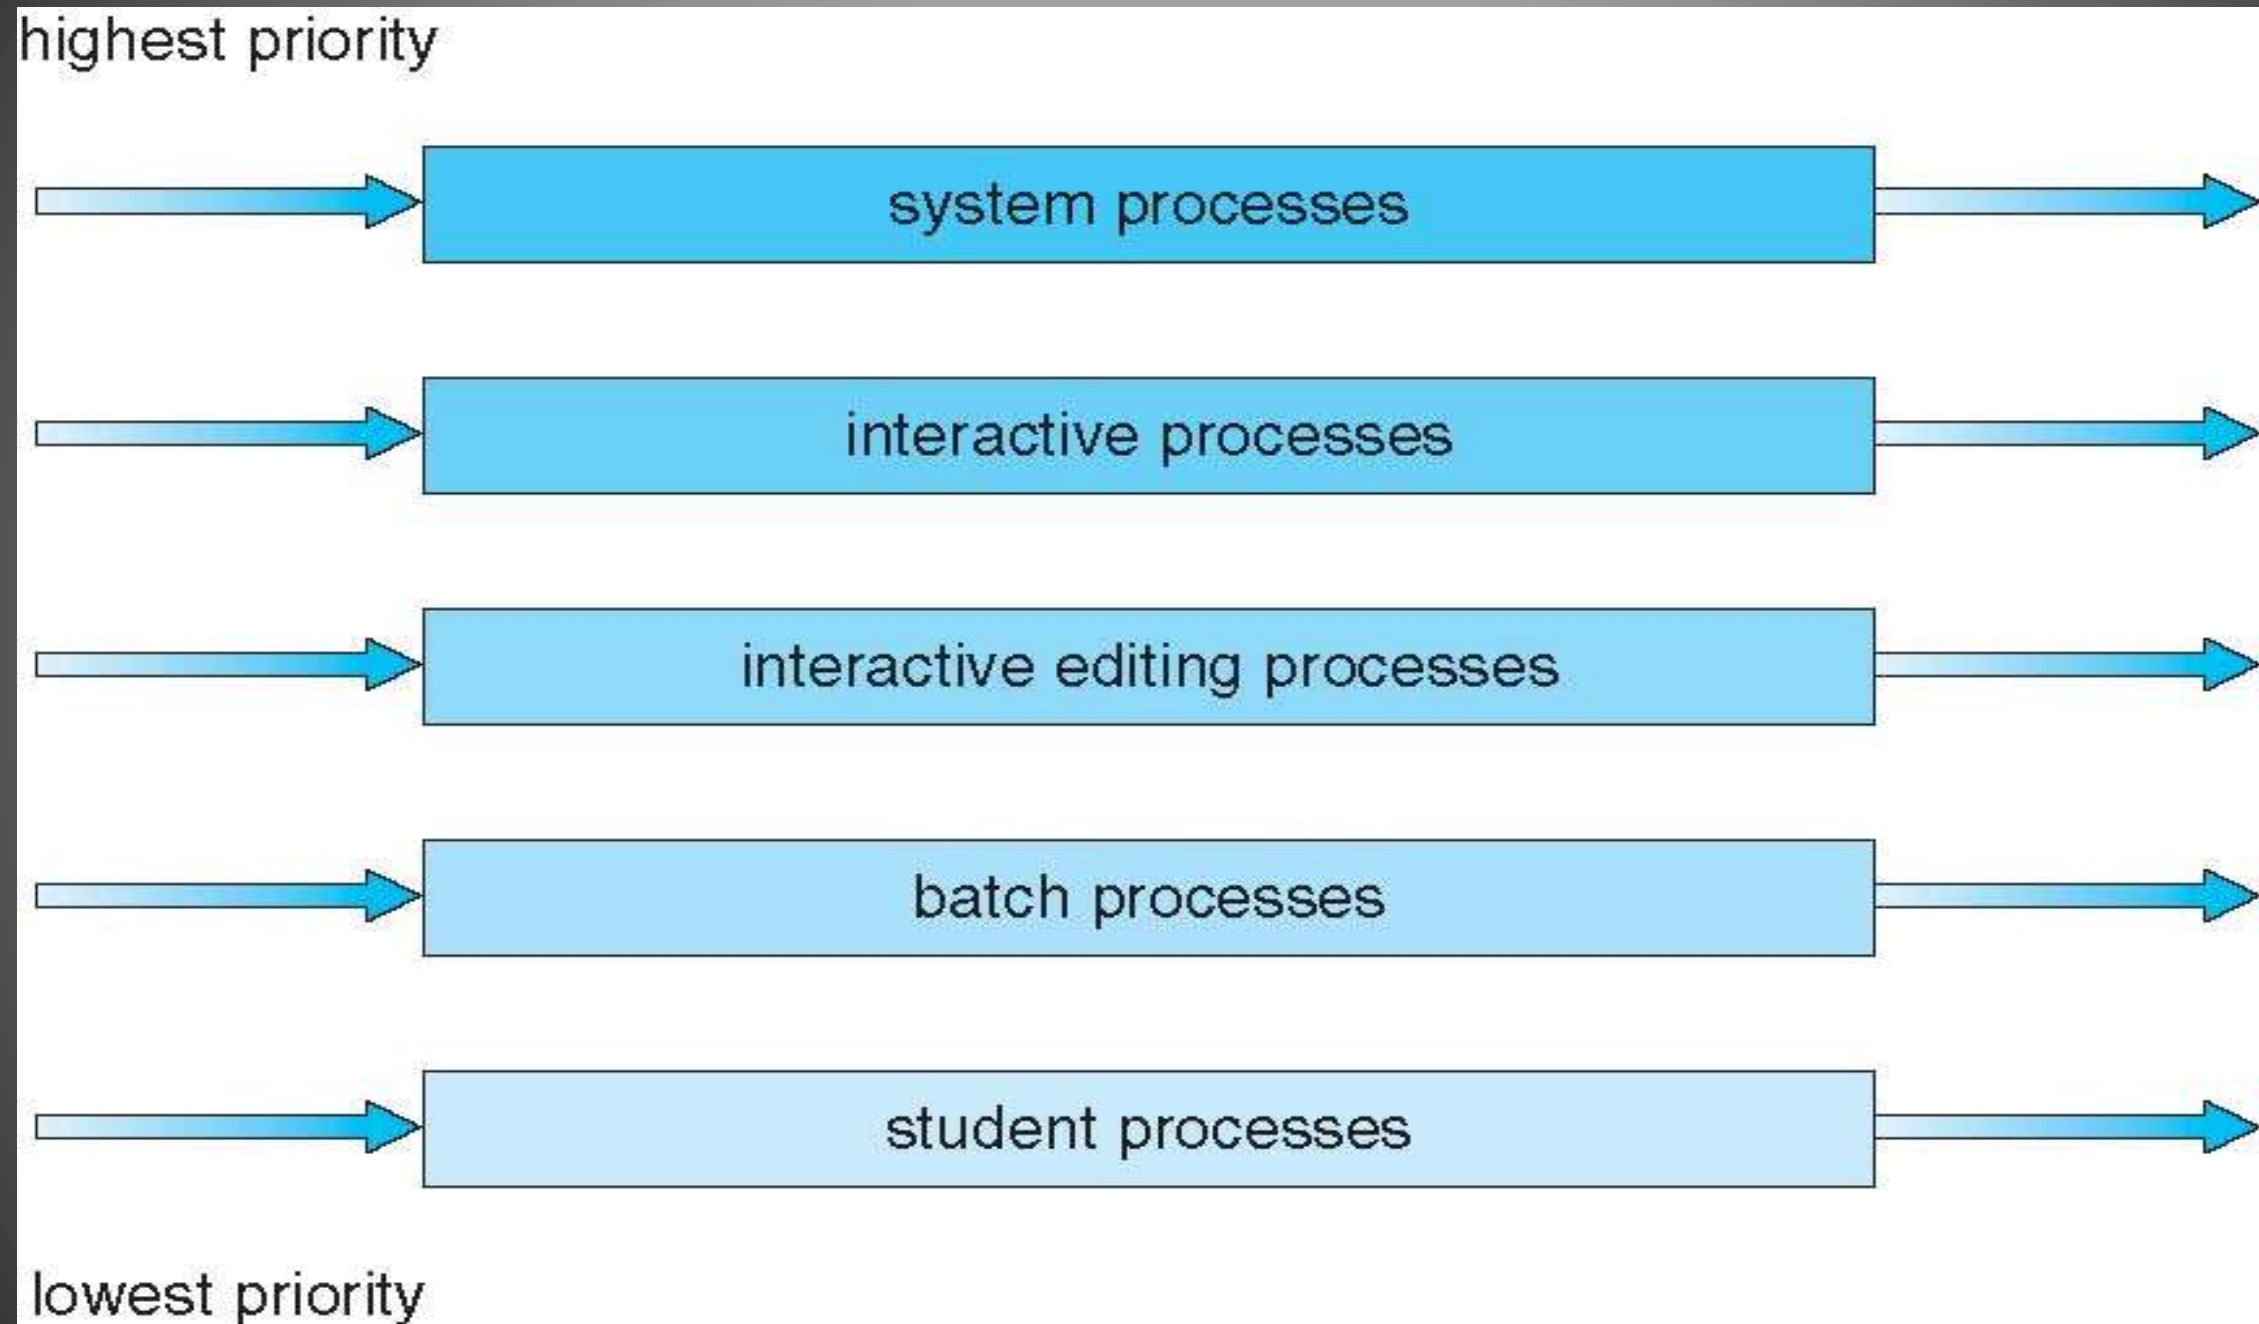
Example RR

- ▶ Covered in class

Multilevel Queue

- ▶ Process is assigned to one queue based on memory size, priority, process type.
- ▶ Ready queue is partitioned into separate queues, eg:
 - **foreground** (interactive)
 - **background** (batch)
- ▶ Each queue has its own scheduling algorithm:
 - foreground – RR
 - background – FCFS
- ▶ Scheduling must be done between the queues:
 - Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
 - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR
 - 20% to background in FCFS

Multilevel Queue Scheduling




Multilevel Feedback Queue


- ▶ Idea is to separate process according to CPU burst time
- ▶ If a process uses too much CPU time, it is moved to low priority
- ▶ If a process waits too long (aging) in low priority then it is moved to high priority
- ▶ Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service

Example of Multilevel Feedback Queue


- ▶ Three queues:
 - Q_0 – RR with time quantum 8 milliseconds
 - Q_1 – RR time quantum 16 milliseconds
 - Q_2 – FCFS
- ▶ Scheduling
 - A new job enters queue Q_0 which is served FCFS
 - When it gains CPU, job receives 8 milliseconds
 - If it does not finish in 8 milliseconds, job is moved to queue Q_1
 - At Q_1 job is again served FCFS and receives 16 additional milliseconds
 - If it still does not complete, it is preempted and moved to queue Q_2




Multiple-Processor Scheduling

- ▶ CPU scheduling more complex when multiple CPUs are available
 - ▶ **Homogeneous processors** within a multiprocessor
 - ▶ **Asymmetric multiprocessing** – only one processor accesses the system data structures, alleviating the need for data sharing
 - ▶ **Symmetric multiprocessing (SMP)** – each processor is self-scheduling, all processes in common ready queue, or each has its own private queue of ready processes
- 

Multiple-Processor Scheduling

- ▶ **Processor affinity** – process has affinity for processor on which it is currently running
 - ▶ **soft affinity:** When an operating system has a policy of attempting to keep a process running on the same processor—but not guaranteeing that it will do so known as **soft affinity**.
 - ▶ **hard affinity:** allowing a process to specify a subset of processors on which it may run.
- 

Multiple-Processor Scheduling – Load Balancing

- ▶ If SMP, need to keep all CPUs loaded for efficiency
 - ▶ **Load balancing** attempts to keep workload evenly distributed
 - ▶ **Push migration** – periodic task checks load on each processor, and if found pushes task from overloaded CPU to other CPUs
 - ▶ **Pull migration** – idle processors pulls waiting task from busy processor
- 

Real-Time CPU Scheduling

- ▶ **Soft real-time systems** – no guarantee as to when critical real-time process will be scheduled
- ▶ **Hard real-time systems** – task must be serviced by its deadline
- ▶ Two types of latencies affect performance
 1. Interrupt latency – time from arrival of interrupt to start of routine that services interrupt
 2. Dispatch latency – time for schedule to take current process off CPU and switch to another

