

Process Synchronization

Course Instructor: Nausheen Shoaib

Multiple Processes

- ▶ Operating Systems is managing multiple processes:
 - Multiprogramming- The management of multiple processes within a uniprocessor system.
 - Multiprocessing- The management of multiple processes within a multiprocessor.
 - Distributed Processing- The management of multiple processes executing on multiple, distributed computer systems.e.g clusters
- ▶ Big Issue is Concurrency

Concurrency [1 / 2]

- ▶ Concurrency encompasses a host of design issues, including :
 - communication among processes
 - sharing of and competing for resources (such as memory, files, and I/O access)
 - synchronization of the activities of multiple processes
 - allocation of processor time to processes

Concurrency [2 / 2]

Concurrency arises in:

- ▶ **Multiple applications**
 - Sharing time
- ▶ **Structured applications**
 - Extension of modular design
- ▶ **Operating system structure**
 - OS themselves implemented as a set of processes or threads

Difficulties of Concurrency

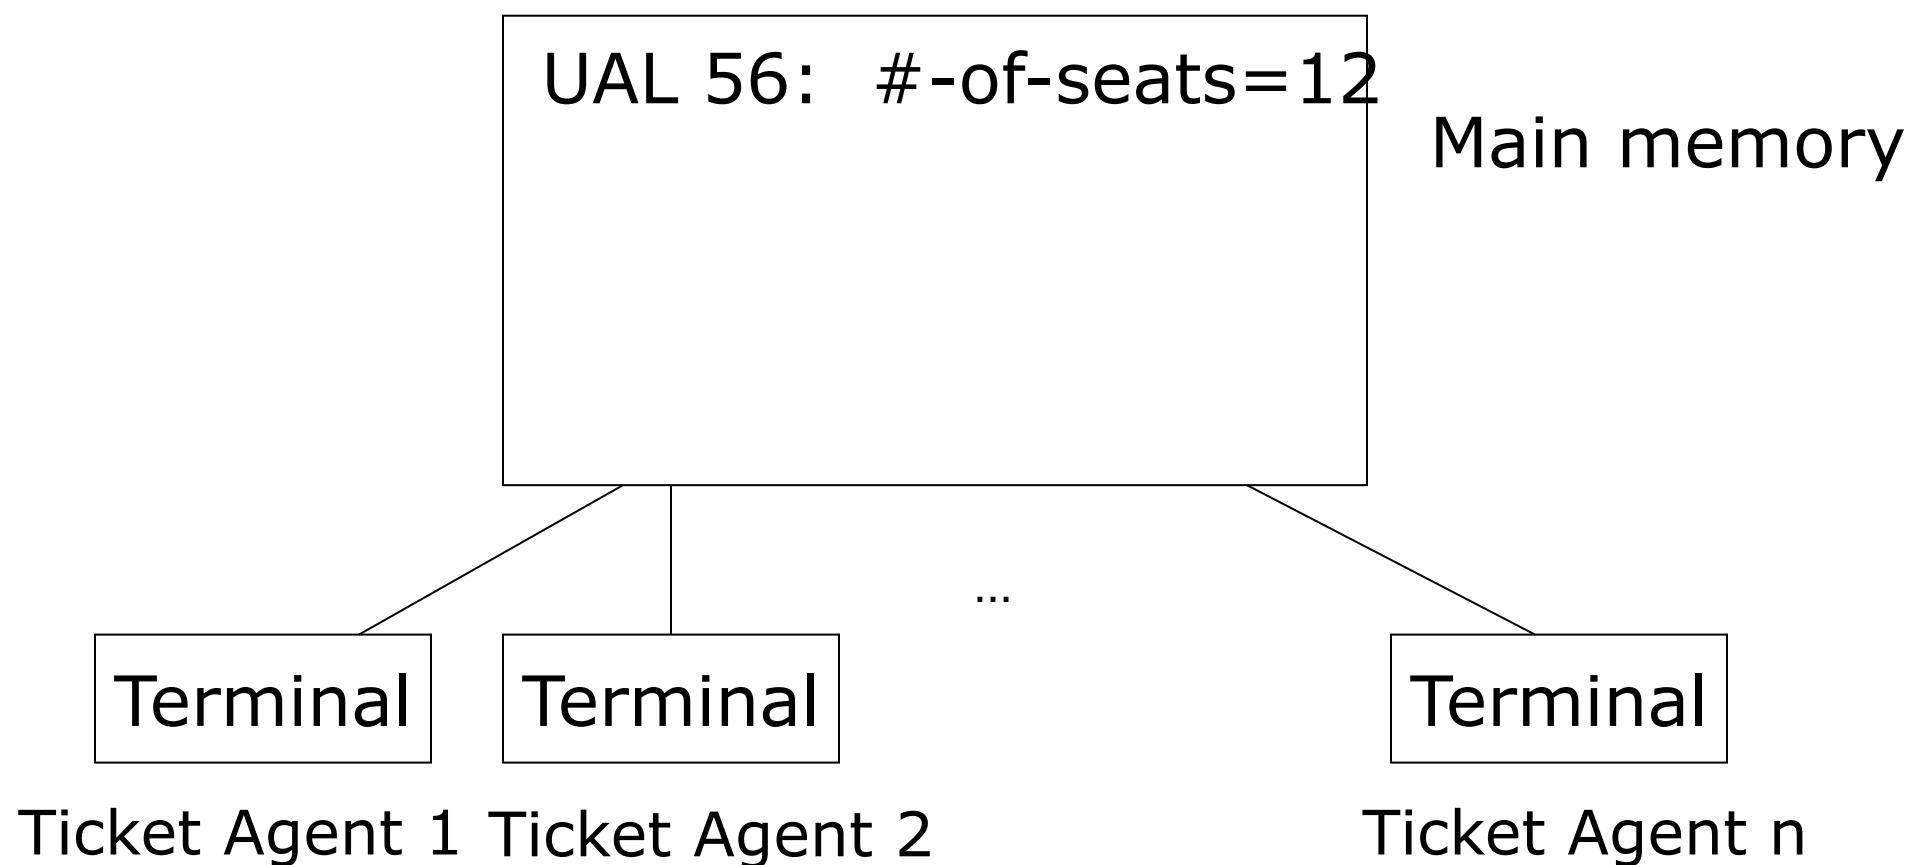
- ▶ Sharing of global resources
 - Writing a shared variable: the order of writes is important
 - Incomplete writes a major problem
- ▶ Optimally managing the allocation of resources
- ▶ Difficult to locate programming errors as results are not deterministic and reproducible.

Race Condition

- ▶ A situation where several processes access and manipulate the same data concurrently and the outcome of the execution depends on the particular order in which the access takes place.

Example for Race condition

- ▶ Suppose a customer wants to book a seat on UAL 56. Ticket agent will check the #‐of‐seats. If it is greater than 0, he will grab a seat and decrement #‐of‐seats by 1.



Example for Race condition(cont.)

Ticket Agent 1

P1: LOAD #-of-seats
P2: DEC 1
P3: STORE #-of-seats

Ticket Agent 2

Q1: LOAD #-of-seats
Q2: DEC 1
Q3: STORE #-of-seats

Ticket Agent 3

R1: LOAD #-of-seats
R2: DEC 1
R3: STORE #-of-seats

Suppose instructions are interleaved as P1,Q1,R1,P2,Q2,R2,P3,Q3,R3

To solve the above problem, we must make sure that:

P1,P2,P3 must be completely executed before we execute Q1 or R1, or
Q1,Q2,Q3 must be completely executed before we execute P1 or R1, or
R1,R2,R3 must be completely executed before we execute P1 or Q1.

The Critical Section / Region Problem

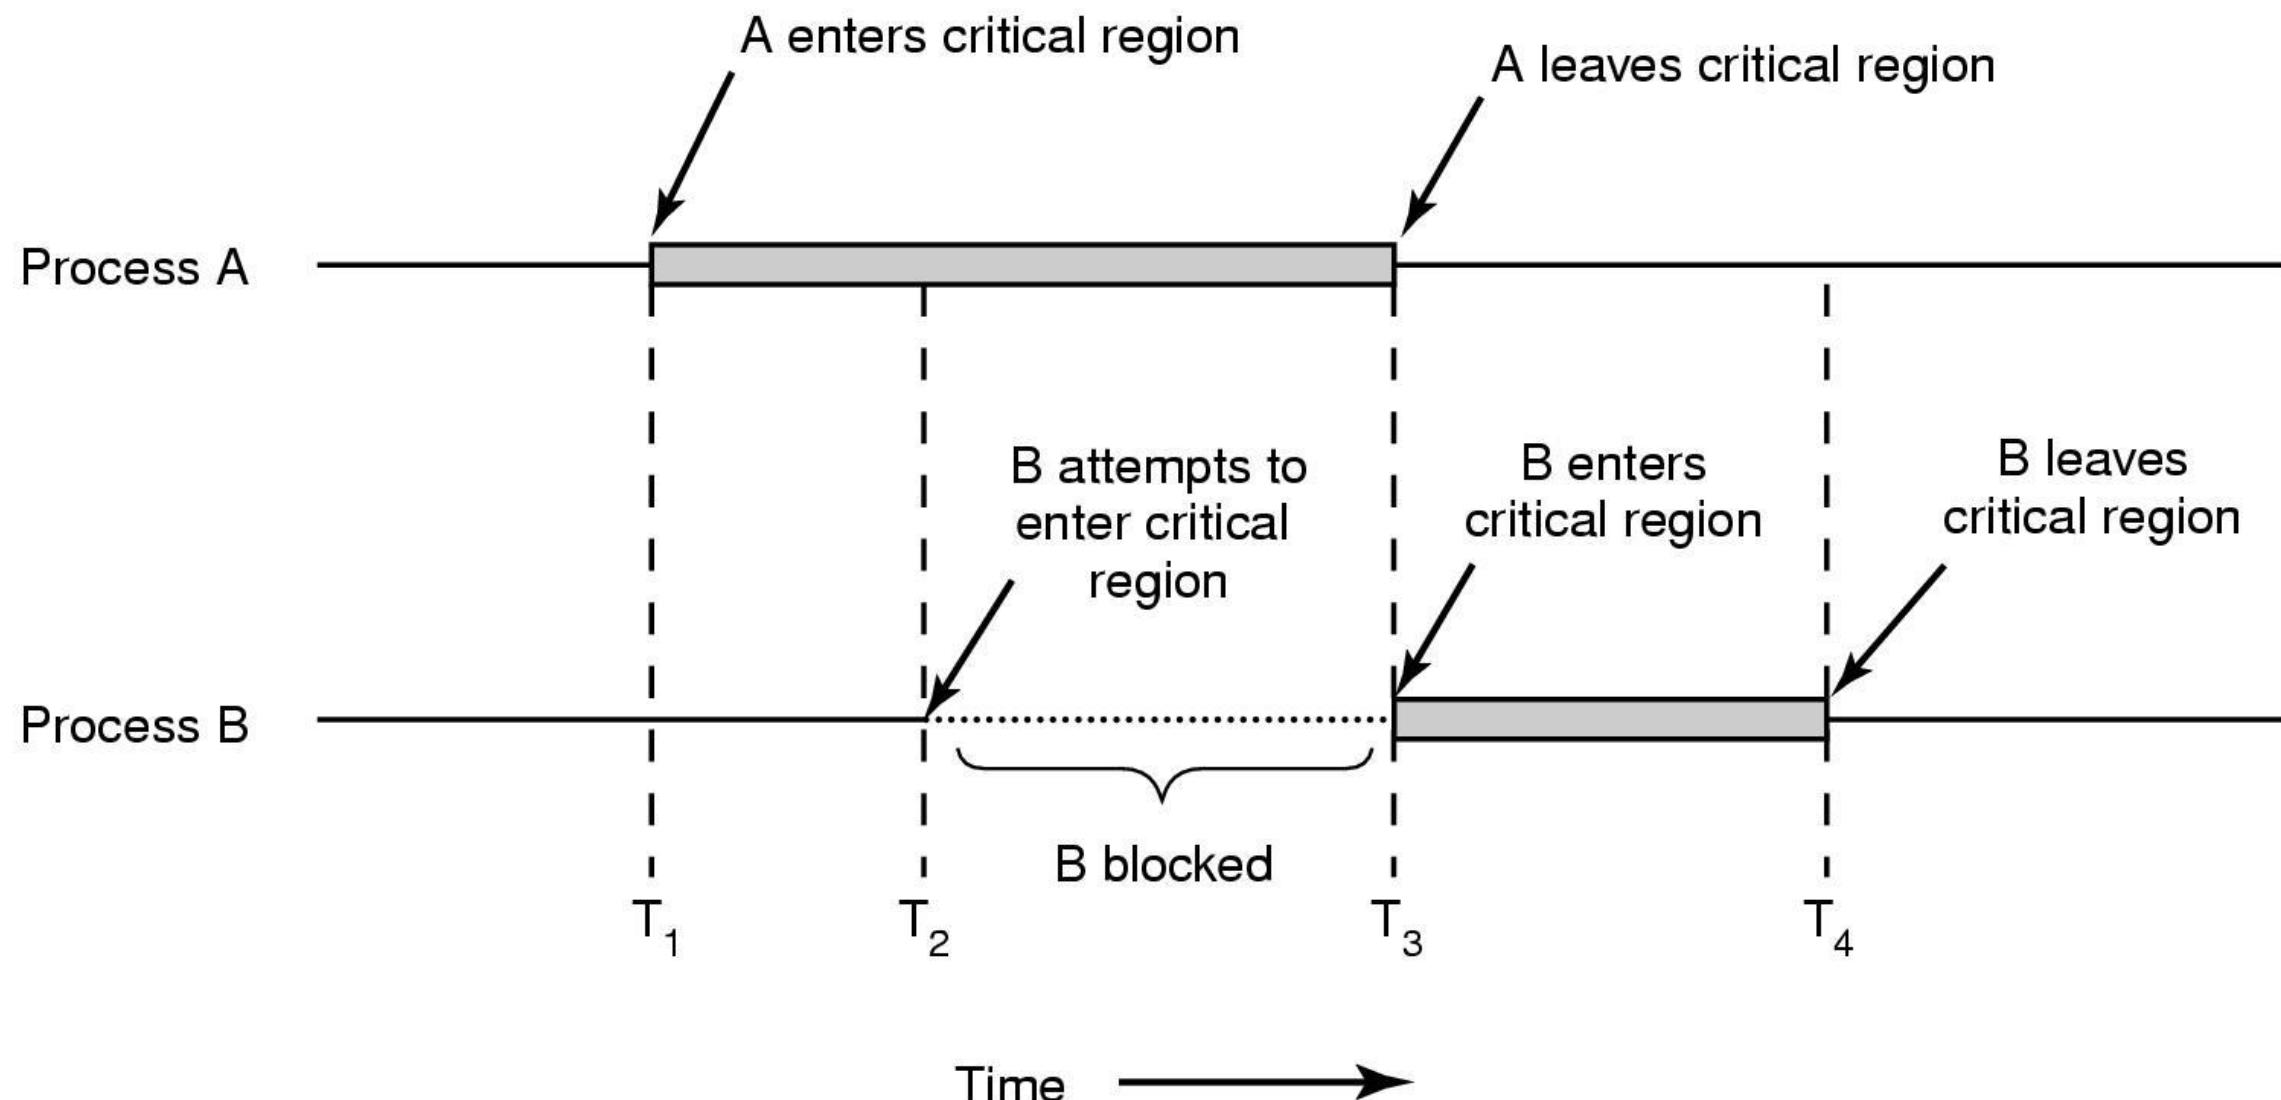
- ▶ Occurs in systems where multiple processes all compete for the use of shared data.
- ▶ Each process includes a section of code (**the critical section**) where it accesses this shared data.
- ▶ The problem is to ensure that **only one process at a time is allowed** to be operating in its critical section.

Critical Regions (1)

Conditions required to avoid race condition:

- No two processes may be simultaneously inside their critical regions.
- No assumptions may be made about speeds or the number of CPUs.
- No process running outside its critical region may block other processes.
- No process should have to wait forever to enter its critical region.

Critical Regions (2)

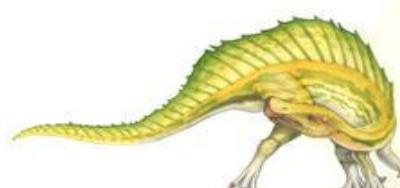


Mutual exclusion using critical regions.



Solution to Critical-Section Problem

1. **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections
2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted
 - Assume that each process executes at a nonzero speed
 - No assumption concerning **relative speed** of the n processes





Critical-Section Handling in OS

Two approaches depending on if kernel is preemptive or non-preemptive

- **Preemptive** – allows preemption of process when running in kernel mode
- **Non-preemptive** – runs until exits kernel mode, blocks, or voluntarily yields CPU
 - ▶ Essentially free of race conditions in kernel mode

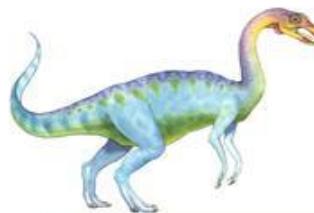




Peterson's Solution

- Good algorithmic description of solving the problem
- Two process solution
- Assume that the `load` and `store` machine-language instructions are atomic; that is, cannot be interrupted
- The two processes share two variables:
 - `int turn;`
 - `Boolean flag[2]`
- The variable `turn` indicates whose turn it is to enter the critical section
- The `flag` array is used to indicate if a process is ready to enter the critical section. `flag[i] = true` implies that process P_i is ready!

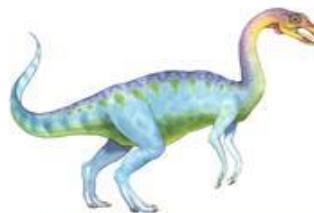




Algorithm for Process P_i

```
do {  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j);  
        critical section  
    flag[i] = false;  
        remainder section  
} while (true);
```





Peterson's Solution (Cont.)

- Provable that the three CS requirement are met:

1. Mutual exclusion is preserved

P_i enters CS only if:

either `flag[j] = false` or `turn = i`

2. Progress requirement is satisfied

3. Bounded-waiting requirement is met



Peterson's Solution and Modern Architecture

- ▶ Although useful for demonstrating an algorithm, Peterson's Solution is not guaranteed to work on modern architectures.
 - To improve performance, processors and/or compilers may reorder operations that have no dependencies
- ▶ Understanding why it will not work is useful for better understanding race conditions.
- ▶ For single-threaded this is ok as the result will always be the same.
- ▶ For multithreaded the reordering may produce inconsistent or unexpected results!

Modern Architecture Example

- ▶ Two threads share the data:

```
boolean flag = false;
```

```
int x = 0;
```

- ▶ Thread 1 performs

```
while (!flag)
```

```
;
```

```
    print x
```

- ▶ Thread 2 performs

```
    x = 100;
```

```
    flag = true
```

- ▶ What is the expected output?

Modern Architecture Example (Cont.)

- ▶ However, since the variables `flag` and `x` are independent of each other, the instructions:

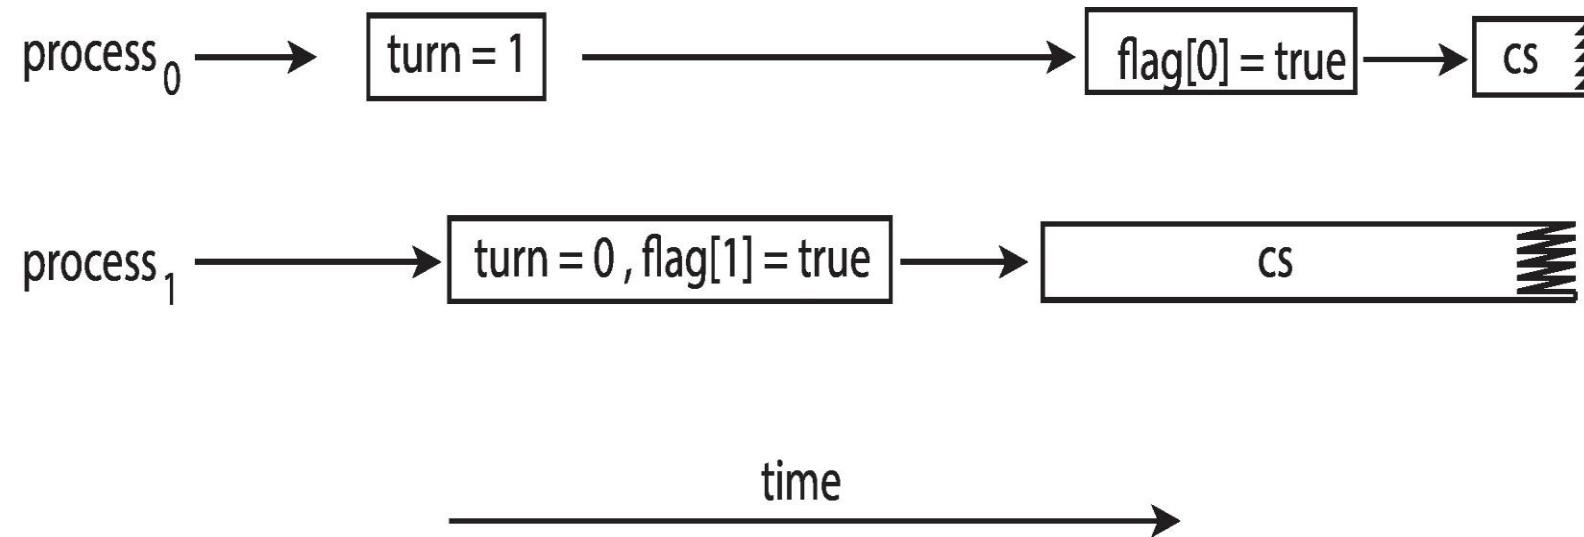
```
flag = true;  
x = 100;
```

for Thread 2 may be reordered

- ▶ If this occurs, the output may be 0!

Peterson's Solution Revisited

- ▶ The effects of instruction reordering in Peterson's Solution



- ▶ This allows both processes to be in their critical section at the same time!
- ▶ To ensure that Peterson's solution will work correctly on modern computer architecture we must use **Memory Barrier**.

Memory Barrier

- ▶ **Memory model** are the memory guarantees a computer architecture makes to application programs.
- ▶ Memory models may be either:
 - **Strongly ordered** – where a memory modification of one processor is immediately visible to all other processors.
 - **Weakly ordered** – where a memory modification of one processor may not be immediately visible to all other processors.
- ▶ A **memory barrier** is an instruction that forces any change in memory to be propagated (made visible) to all other processors.

Memory Barrier Instructions

- ▶ When a memory barrier instruction is performed, the system ensures that all loads and stores are completed before any subsequent load or store operations are performed.
- ▶ Therefore, even if instructions were reordered, the memory barrier ensures that the store operations are completed in memory and visible to other processors before future load or store operations are performed.

Synchronization Hardware

- ▶ Many systems provide hardware support for implementing the critical section code.
- ▶ Uniprocessors – could disable interrupts
 - Currently running code would execute without preemption
 - Generally too inefficient on multiprocessor systems
 - Operating systems using this not broadly scalable
- ▶ We will look at three forms of hardware support:
 1. Hardware instructions
 2. Atomic variables

Hardware Instructions

- ▶ Special hardware instructions that allow us to either *test-and-modify* the content of a word, or two *swap* the contents of two words atomically (uninterruptedly.)
 - Test-and-Set instruction
 - Compare-and-Swap instruction

The test_and_set Instruction

- ▶ Definition

```
boolean test_and_set  
(boolean *target)  
{  
    boolean rv =  
*target;  
    *target = true;  
    return rv;  
}
```

The test_and_set Instruction

- ▶ if two test and set() instructions are executed simultaneously(each on a different core), they will be executed sequentially in some arbitrary order.
- ▶ If the machine supports the test and set() instruction, then we can implement mutual exclusion by declaring a boolean variable lock, initialized to false.

Solution Using test_and_set()

- ▶ Shared boolean variable **lock**, initialized to **false**
- ▶ Solution:

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = false;  
    /* remainder section */  
}  
while (true);
```

The compare_and_swap Instruction

- ▶ Definition

```
int compare_and_swap(int *value, int expected, int new_value)
```

```
{  
    int temp = *value;  
    if (*value == expected)  
        *value = new_value;  
    return temp;
```

```
}
```

- ▶ Properties

- Executed atomically
- Returns the original value of passed parameter **value**
- Set the variable **value** the value of the passed parameter **new_value** but only if ***value == expected** is true. That is, the swap takes place only under this condition.

The compare_and_swap Instruction

- ▶ if two CAS instructions are executed simultaneously (each on a different core), they will be executed sequentially in some arbitrary order.
- ▶ Mutual exclusion using CAS can be provided

Solution using compare_and_swap

- ▶ Shared integer **lock** initialized to 0;
- ▶ Solution:

```
while (true) {
    while (compare_and_swap(&lock, 0, 1)
!= 0)
        ; /* do nothing */

    /* critical section */

    lock = 0;

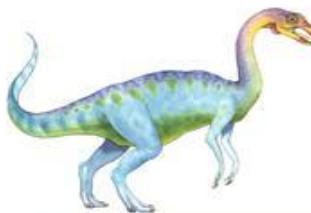
    /* remainder section */
}
```

Atomic Variables

- ▶ Typically, instructions such as compare-and-swap are used as building blocks for other synchronization tools.
- ▶ One tool is an **atomic variable** that provides *atomic* (uninterruptible) updates on basic data types such as integers and booleans.

Mutex Locks

- n OS designers build software tools to solve critical section problem
- n Simplest is mutex lock
- n Protect critical regions with it by first **acquire()** a lock then **release()** it
 - | Boolean variable indicating if lock is available or not
- n Calls to **acquire()** and **release()** must be atomic
 - | Usually implemented via hardware atomic instructions
- n But this solution requires **busy waiting**
- n **Busy Waiting means** busy-looping or spinning is a technique in which a process repeatedly checks to see if a condition is true, such as whether keyboard input or a lock is available
- n Mutex lock therefore called a **spinlock**



acquire() and release()

- `acquire() {`
 `while (!available)`
 `; /* busy wait */`
 `available = false;`
}
- `release() {`
 `available = true;`
}
- `do {`
 `acquire lock`
 `critical section`
 `release lock`
 `remainder section`
`} while (true);`



Pthreads Synchronization – Mutex Lock

```
#include <pthread.h>
Pthread_mutex_t mutex;

/* create the mutex lock */
Pthread_mutex_init(&mutex, NULL);
```

1st Arg: Mutex initilizer

2nd Arg: Attributes, NULL means no error checks will be performed

Pthreads Synchronization – Mutex Lock

- ▶ The mutex is acquired and released with the `pthread_mutex_lock()`
- ▶ and `pthread_mutex_unlock()` functions.
- ▶ If the mutex lock is unavailable when `pthread_mutex_lock()` is invoked, the calling thread is blocked until the owner invokes `pthread_mutex_unlock()`.

What is a Semaphore

- ▶ a **semaphore** is a variable or abstract data type used to control access to a common resource by multiple processes in a concurrent system such as a multiprogramming operating system.

Semaphore

- ▶ Synchronization tool that does not require busy waiting
- ▶ Semaphore S – integer variable
- ▶ Two standard operations modify S : `wait()` and `signal()`
- ▶ Less complicated
- ▶ Can only be accessed via two indivisible (atomic) operations

Definition of `wait()` operation:

```
wait (S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

Definition of `signal()` operation:

```
signal (S) {  
    S++;  
}
```

- ▶ when one process modifies the semaphore value, no other process can simultaneously modify that same semaphore value.

Semaphore Usage

Counting Semaphore:

- ▶ Semaphores which allow an arbitrary resource count are called **counting semaphores**,
- ▶ Counting semaphores can be used to control access to a given resource consisting of a finite number of instances
- ▶ Can implement a counting semaphore S as a binary semaphore

Binary Semaphore:

- ▶ while semaphores which are restricted to the values 0 and 1 (or locked/unlocked, unavailable/available) are called **binary semaphores** and are used to implement locks.
- ▶ Binary semaphores behave similarly to mutex locks.
- ▶ Can implement a counting semaphore S as a binary semaphore

Semaphore Usage

- ▶ consider two concurrently running processes: P_1 with a statement S_1 and P_2 with a statement S_2 .
- ▶ It is required that S_2 be executed only after S_1 has completed. We can implement this scheme readily by letting P_1 and P_2 share a common semaphore synch , initialized to 0.

P_1 :

```
 $S_1;$   
signal(synch);
```

P_2 :

```
wait(synch);  
 $S_2;$ 
```

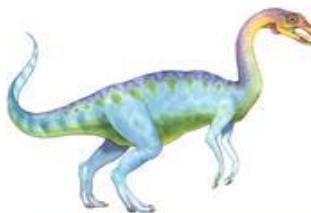
- ▶ Because synch is initialized to 0, P_2 will execute S_2 only after P_1 has invoked $\text{signal}(\text{synch})$, which is after statement S_1 has been executed.

Semaphore Implementation

- ▶ Must guarantee that no two processes can execute `wait()` and `signal()` on the same semaphore at the same time
- ▶ Thus, implementation becomes the critical section problem where the wait and signal code are placed in the critical section
 - Could now have **busy waiting** in critical section implementation
- ▶ Note that applications may spend lots of time in critical sections and therefore this is not a good solution

Semaphore Implementation with no Busy waiting

- ▶ With each semaphore there is an associated waiting queue
- ▶ Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue



Implementation with no Busy waiting (Cont.)

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}  
  
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```



POSIX – Semaphores

```
#include <semaphore.h>
Sem_t sem;
/* Create the semaphore and initialize it to 1 */
```

```
Sem_init(&sem, 0, 1);
```

The `sem_init()` function is passed three parameters:

1. A pointer to the semaphore
2. A flag indicating the level of sharing
3. The semaphore's initial value

POSIX – Semaphores

```
/* acquire the semaphore */
```

```
Sem_wait(&sem);
```

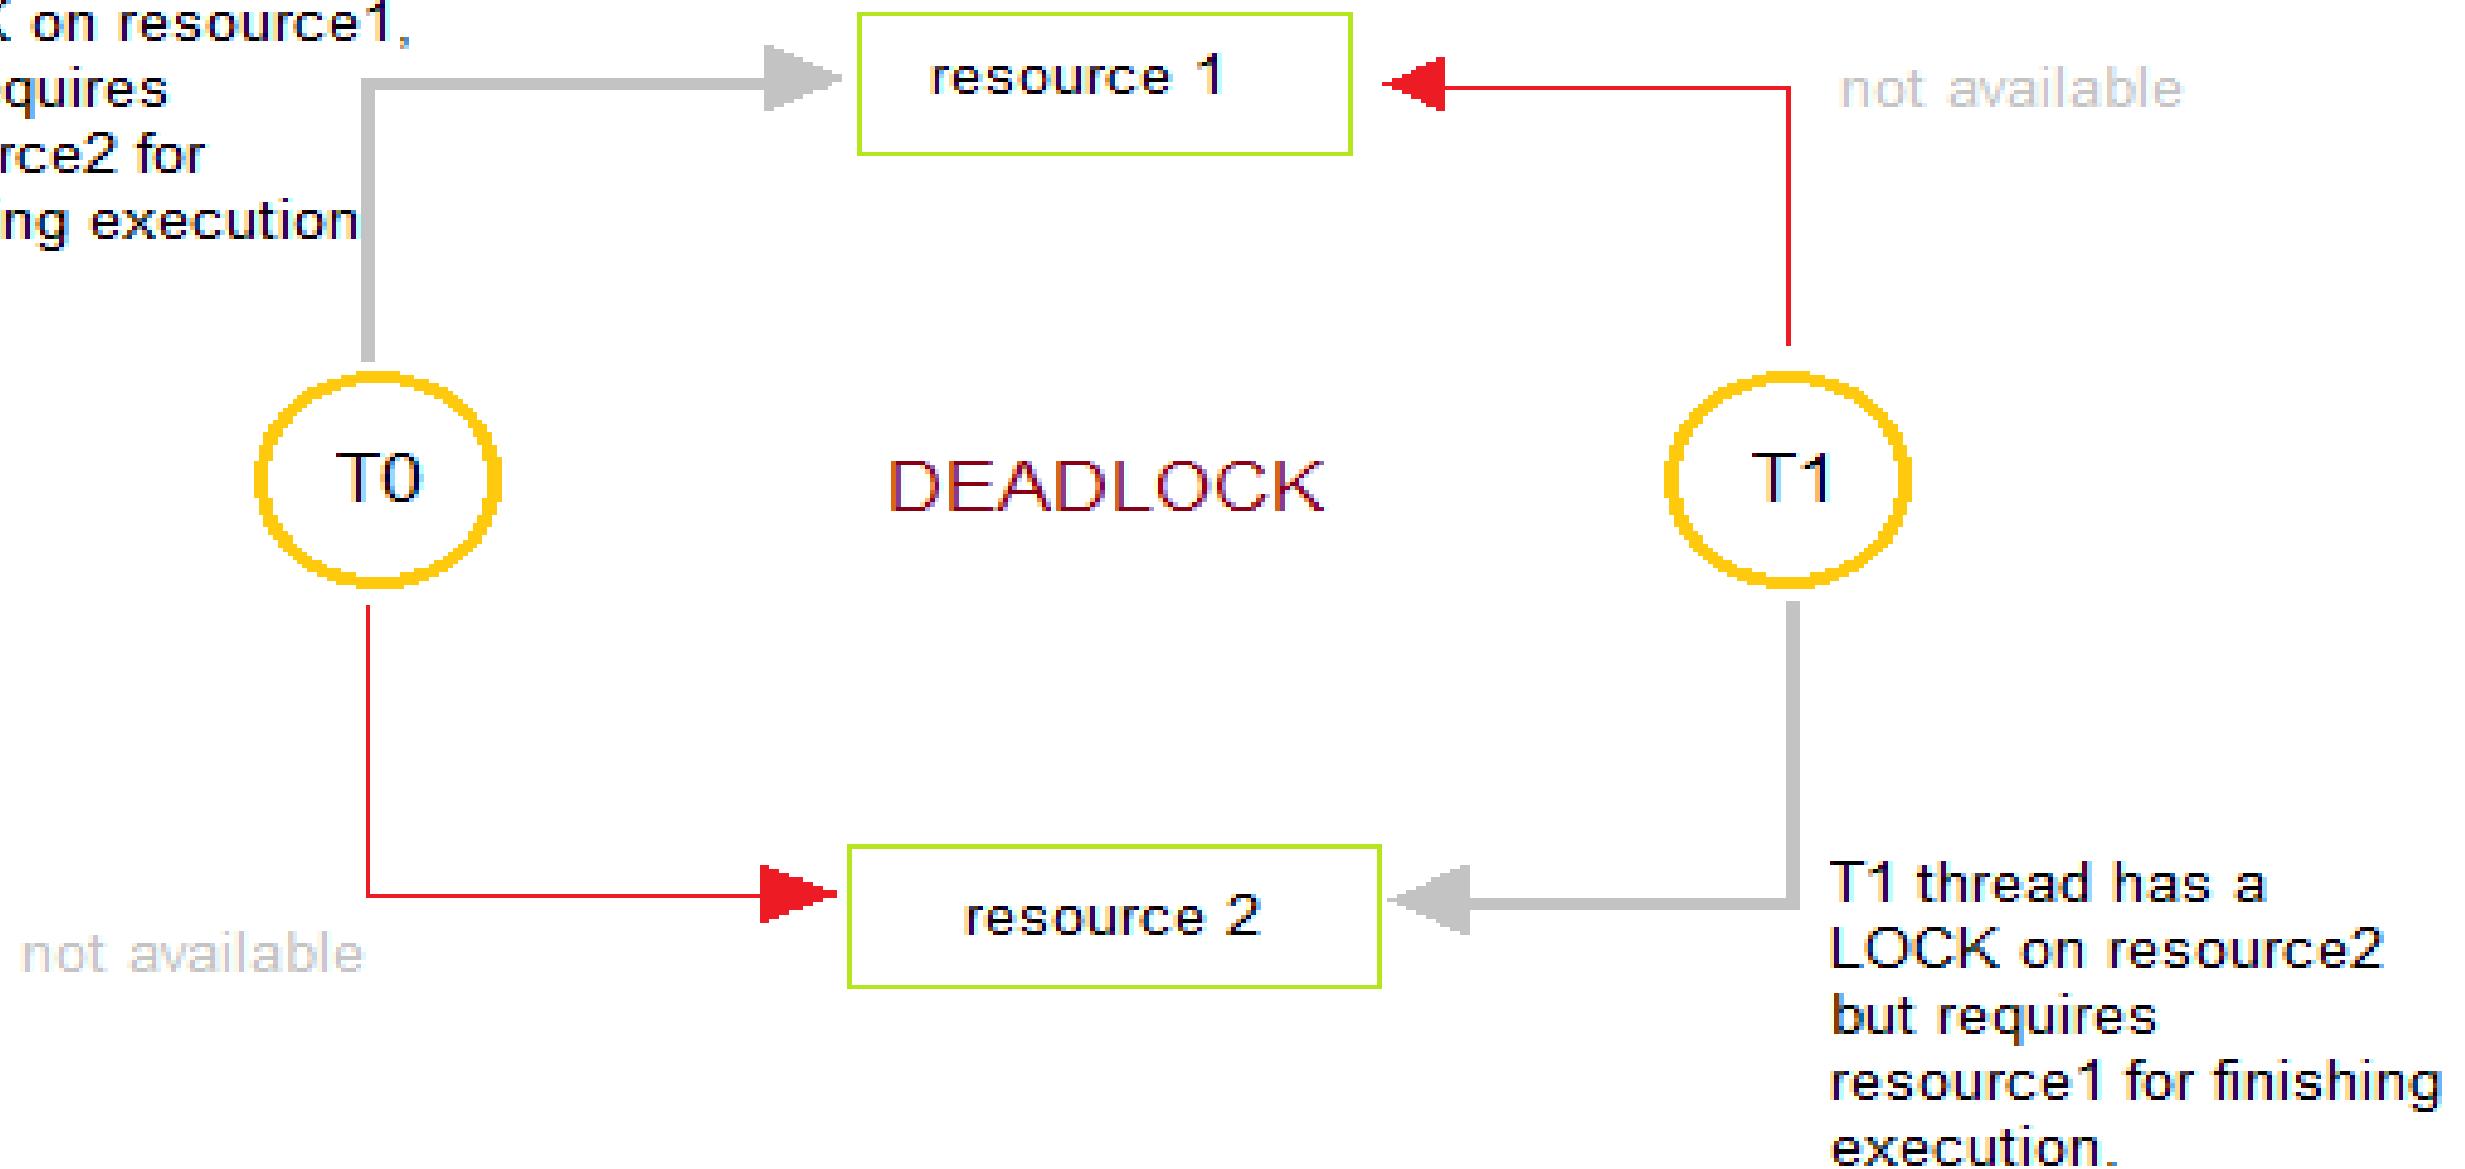
```
/* critical section */
```

```
/* release the semaphore */
```

```
Sem_post(&sem);
```

Deadlock

T0 thread has a
LOCK on resource1,
but requires
resource2 for
finishing execution



Starvation

- ▶ **Starvation** is the name given to the indefinite postponement of a process because it requires some resource before it can run, but the resource, though available for allocation, is never allocated to this process.

Deadlock Vs. Starvation

- ▶ Deadlock refers to the situation when processes are stuck in circular waiting for the resources.
- ▶ On the other hand, starvation occurs when a process waits for a resource indefinitely.

Priority Inversion

- ▶ **Priority Inversion** – Scheduling problem when lower-priority process holds a lock needed by higher-priority process
 - Solved via **priority-inheritance protocol**

Priority Inheritance Protocol

- ▶ when a job blocks one or more high-priority jobs, it ignores its original priority assignment and executes its critical section at an elevated priority level.
- ▶ After executing its critical section and releasing its locks, the process returns to its original priority level

Classical Problems of Synchronization

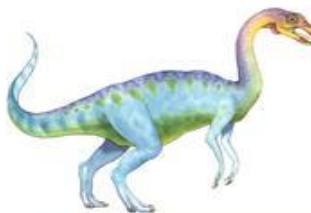
- ▶ Classical problems used to test newly-proposed synchronization schemes
 - Bounded Buffer Problem
 - Dining–Philosophers Problem
 - Readers and Writers Problem



Bounded-Buffer Problem

- **n** buffers, each can hold one item
- Semaphore **mutex** initialized to the value 1
- Semaphore **full** initialized to the value 0
- Semaphore **empty** initialized to the value n





Bounded Buffer Problem (Cont.)

- The structure of the producer process

```
do {  
    ...  
    /* produce an item in next_produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    /* add next produced to the buffer */  
    ...  
    signal(mutex);  
    signal(full);  
} while (true);
```





Bounded Buffer Problem (Cont.)

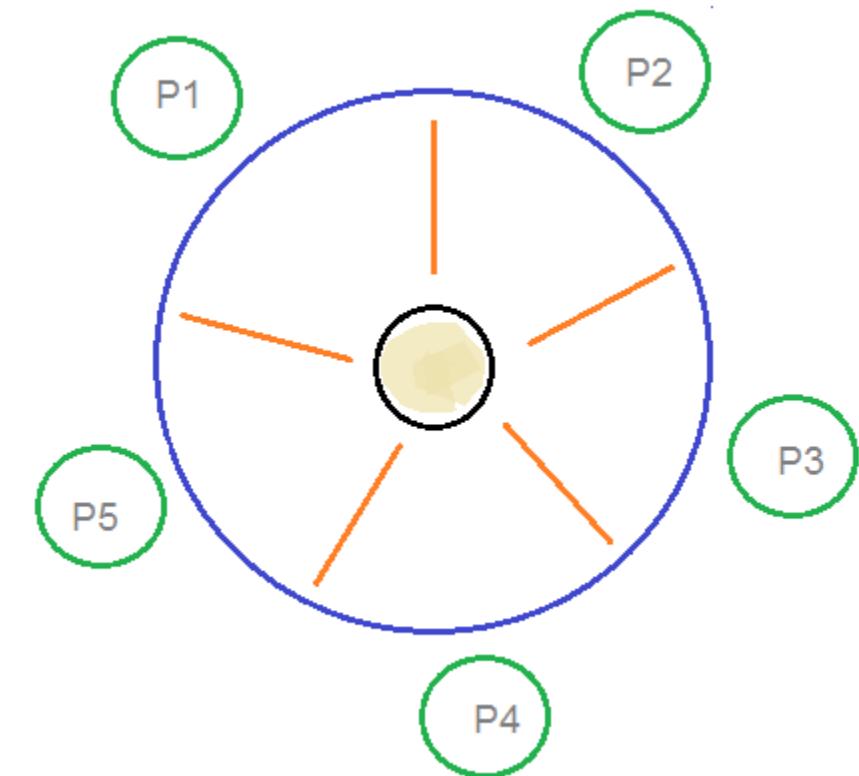
- The structure of the consumer process

```
Do {  
    wait(full);  
    wait(mutex);  
    ...  
    /* remove an item from buffer to next_consumed */  
    ...  
    signal(mutex);  
    signal(empty);  
    ...  
    /* consume the item in next consumed */  
    ...  
} while (true);
```



Dining-Philosophers Problem

- Consider there are five philosophers sitting around a circular dining table. The dining table has five chopsticks and a bowl of rice.



Dining-Philosophers Solution

- ▶ it is clear that a philosopher can think for an indefinite amount of time. But when a philosopher starts eating, he has to stop at some point of time

Dining-Philosophers Solution 1

- ▶ An array of five semaphores, **stick[5]**, for each of the five chopsticks.
- ▶ The code for each philosopher looks like:

```
while(TRUE) {  
    wait(stick[i]);  
    wait(stick[(i+1) % 5]); // mod is used because if i=5, next  
                           // chopstick is 1 (dining table is circular)  
    /* eat */  
    signal(stick[i]);  
    signal(stick[(i+1) % 5]);  
    /* think */  
}
```

Dining-Philosophers Solution 1

- ▶ When a philosopher wants to eat the rice, he will wait for the chopstick at his left and picks up that chopstick.
- ▶ Then he waits for the right chopstick to be available, and then picks it too. After eating, he puts both the chopsticks down

Dining-Philosophers Problem

- ▶ But if all five philosophers are hungry simultaneously, and each of them pickup one chopstick, then a deadlock situation occurs:

Deadlock implies starvation but starvation does not imply deadlock

Dining-Philosophers Solution 2

- ▶ Two Possible solutions are :
- ▶ A philosopher must be allowed to pick up the chopsticks only if both the left and right chopsticks are available.
- ▶ Allow only four philosophers to sit at the table. That way, if all the four philosophers pick up four chopsticks, there will be one chopstick left on the table. So, one philosopher can start eating and eventually, two chopsticks will be available. In this way, deadlocks can be avoided.



Readers-Writers Problem

- A data set is shared among a number of concurrent processes
 - Readers – only read the data set; they do **not** perform any updates
 - Writers – can both read and write
- Problem – allow multiple readers to read at the same time
 - Only one single writer can access the shared data at the same time
- Several variations of how readers and writers are considered – all involve some form of priorities
- Shared Data
 - Data set
 - Semaphore `rw_mutex` initialized to 1
 - Semaphore `mutex` initialized to 1
 - Integer `read_count` initialized to 0



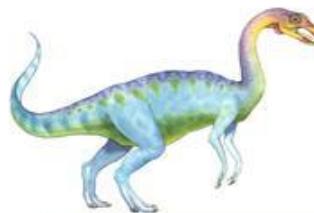


Readers-Writers Problem (Cont.)

- The structure of a writer process

```
do {
    wait(rw_mutex);
    ...
    /* writing is performed */
    ...
    signal(rw_mutex);
} while (true);
```





Readers-Writers Problem (Cont.)

- The structure of a reader process

```
do {  
    wait(mutex);  
    read_count++;  
    if (read_count == 1)  
        wait(rw_mutex);  
    signal(mutex);  
    . . .  
    /* reading is performed */  
    . . .  
    wait(mutex);  
    read_count--;  
    if (read_count == 0)  
        signal(rw_mutex);  
    signal(mutex);  
} while (true);
```

