


Deadlock

Course Instructor: Nausheen Shoaib



Deadlock

- ▶ A set of processes is deadlocked when each process in the set is blocked awaiting an event that can only be triggered by another blocked process in the set
 - Typically involves processes competing for the same set of resources
- ▶ No efficient solution




Potential Deadlock

I need
quad C
and B

I need
quad B
and C

I need
quad D
and A

I need
quad A
and B




Actual Deadlock

HALT until
D is free

HALT until
C is free

HALT until
A is free

HALT until
B is free



Resource Allocation Graphs

- ▶ Directed graph that depicts a state of the system of resources and processes




(a) Resource is requested




(b) Resource is held



Resource Allocation Graphs of deadlock



(c) Circular wait



(d) No deadlock







Resource Allocation Graphs of deadlock

- ▶ Deadlocks can be described more precisely in terms of a directed graph called a **system resource-allocation graph**







Dealing with Deadlock

- ▶ Three general approaches exist for dealing with deadlock.
 - Prevent deadlock
 - Avoid deadlock
 - Detect Deadlock





Deadlock Prevention Strategy

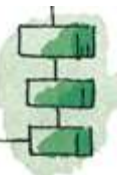
- ▶ Design a system in such a way that the possibility of deadlock is excluded.
- ▶ Two main methods
 - Indirect – prevent one of the three necessary conditions from occurring
 - Direct – prevent circular waits





Deadlock Prevention Conditions

- ▶ Mutual Exclusion
- ▶ Hold and Wait
- ▶ No Preemption
- ▶ Circular Wait





Deadlock Prevention Conditions

Mutual Exclusion

- This condition says, “There exist resources in the system that can be used by only one process at a time.”
- Examples include printer, write access to a file or record, entry into a section of code
- Best not to get rid of this condition
 - some resources are intrinsically nonsharable





Deadlock Prevention Conditions

Hold and Wait (1/2)

- This condition says, "Some process holds one resource while waiting for another."
- To attack the hold and wait condition:
 - Force a process to acquire all the resources it needs before it does anything; if it can't get them all, get none
- Each philosopher tries to get both chopsticks, but if only one is available, put it down and try again later





Deadlock Prevention Conditions

No Preemption (1/2)

- This condition says, "Once a process has a resource, it will not be forced to give it up."
- To attack the no preemption condition:
 - If a process asks for a resource not currently available, block it but also take away all of its other resources
 - Add the preempted resources to the list of resource the blocked process is waiting for






Deadlock Prevention Conditions

Circular Wait (1/2)


- This condition says, “A is blocked waiting for B, B for C, C for D, and D for A”
- Note that the number of processes is actually arbitrary
- To attack the circular wait condition:
 - Assign each resource a priority
 - Make processes acquire resources in priority order





Deadlock Avoidance


- ▶ A decision is made dynamically whether the current resource allocation request will, if granted, potentially lead to a deadlock
- ▶ Requires knowledge of future process requests





Two Approaches to Deadlock Avoidance

- ▶ Process Initiation Denial
- ▶ Resource Allocation Denial





Process Initiation Denial


- ▶ A process is only started if the maximum claim of all current processes plus those of the new process can be met.
- ▶ Not optimal,
 - Assumes the worst: that all processes will make their maximum claims together.






Resource

Allocation Denial



- ▶ Referred to as the banker's algorithm
 - A strategy of resource allocation denial
- ▶ Consider a system with fixed number of resources
 - ***State*** of the system is the current allocation of resources to process
 - ***Safe state*** is where there is at least one sequence that does not result in deadlock
 - ***Unsafe state*** is a state that is not safe




Basic Facts for deadlock avoidance

- ▶ If a system is in safe state \Rightarrow no deadlocks
- ▶ If a system is in unsafe state \Rightarrow possibility of deadlock
- ▶ Avoidance \Rightarrow ensure that a system will never enter an unsafe state.



Banker's Algorithm

- ▶ Multiple instances
- ▶ Each process must a priori claim maximum use
- ▶ When a process requests a resource it may have to wait
- ▶ When a process gets all its resources it must return them in a finite amount of time






Data Structures for the Banker's Algorithm

Let n = number of processes, and m = number of resources types.

- ▶ **Available:** Vector of length m . If available $[j] = k$, there are k instances of resource type R_j available
- ▶ **Max:** $n \times m$ matrix. If $Max[i, j] = k$, then process P_i may request at most k instances of resource type R_j
- ▶ **Allocation:** $n \times m$ matrix. If $Allocation[i, j] = k$ then P_i is currently allocated k instances of R_j
- ▶ **Need:** $n \times m$ matrix. If $Need[i, j] = k$, then P_i may need k more instances of R_j to complete its task


$$Need[i, j] = Max[i, j] - Allocation[i, j]$$





Example of Banker's Algorithm

- ▶ Discussed in Class



An illustration showing a blue box (representing a process) connected by lines to a yellow box (representing a resource) and a red box (representing another resource).

Resource-Request Algorithm for Process P_i

$Request_i$ = request vector for process P_i .
If $Request_i[j] = k$ then process P_i wants k instances of resource type R_j


1. If $Request_i \leq Need_i$, go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
2. If $Request_i \leq Available$, go to step 3. Otherwise P_i must wait, since resources are not available
3. Pretend to allocate requested resources to P_i by modifying the state as follows:

$$Available = Available - Request_i;$$

$$Allocation_i = Allocation_i + Request_i;$$

$$Need_i = Need_i - Request_i;$$

- If safe \Rightarrow the resources are allocated to P_i
- If unsafe $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored





Example of Resource Request Algorithm

- ▶ Discussed in Class

