

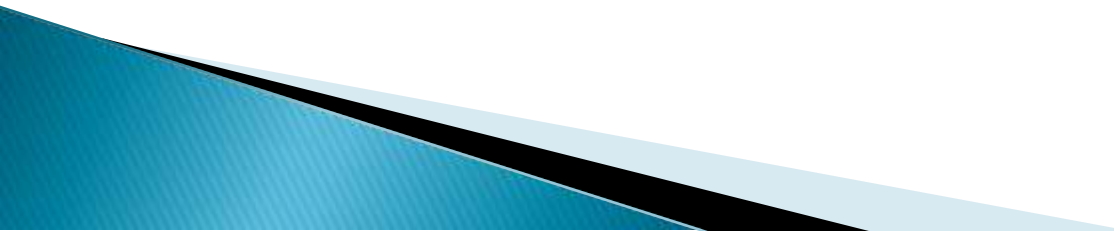
Introduction

Chapter #1

Course Instructor: Nausheen Shoaib



Marks Distribution

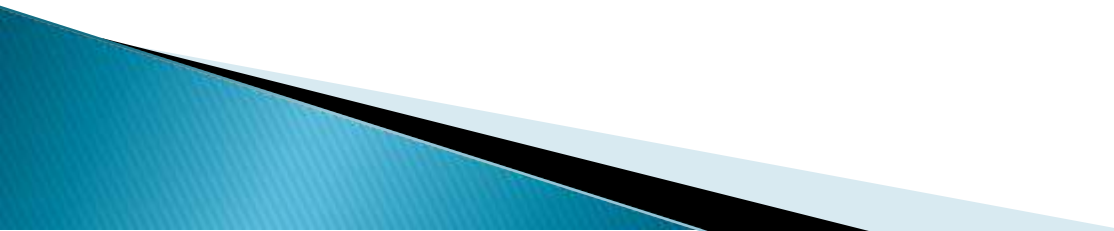
- ▶ Mid1: 15%
 - ▶ Mid2: 15%
 - ▶ Class activities+ Assignment+Projects: 20%
 - ▶ Final: 50%
 - ▶ **Book:** *Operating System Concepts by Abraham Silberschatz 10th Edition*
- 

What is an Operating System?

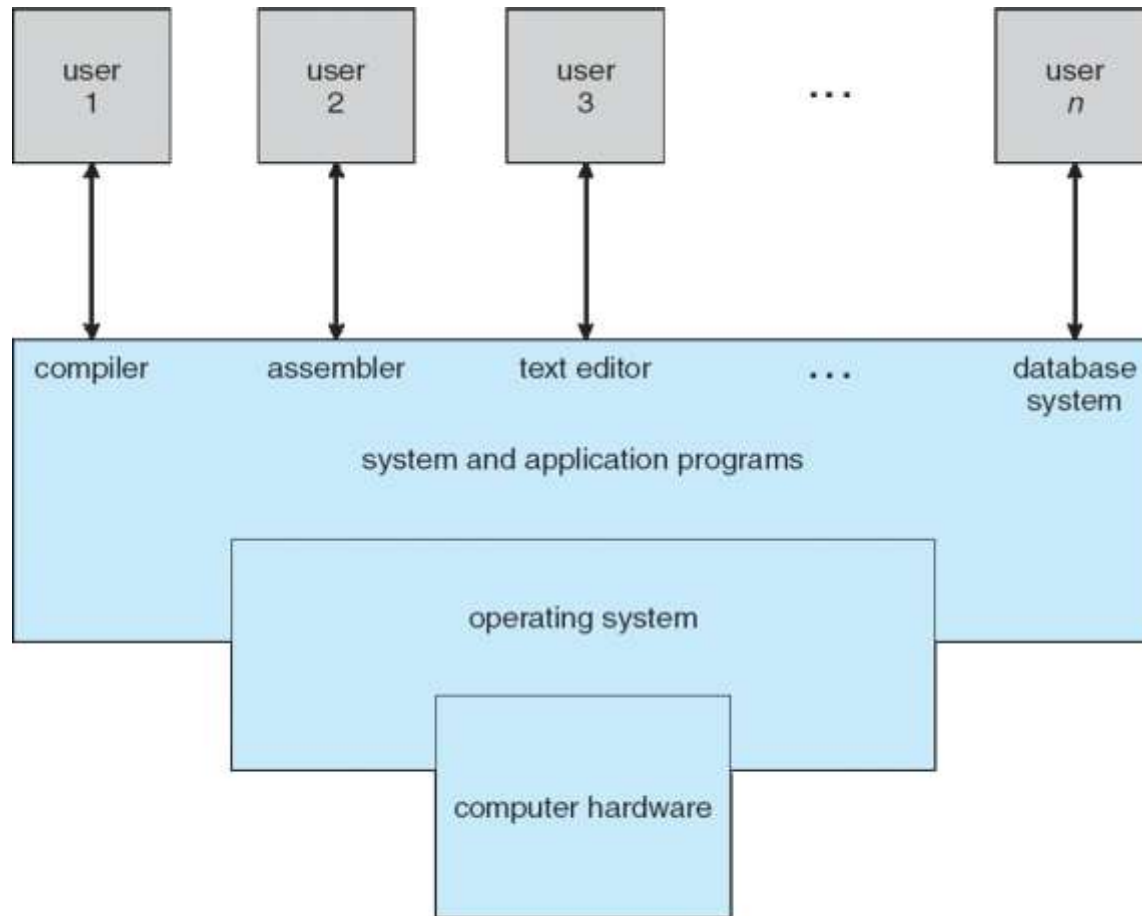
- ▶ A program that acts as an intermediary between a user of a computer and the computer hardware



Computer System Structure

- ▶ Computer system can be divided into four components:
 - Hardware – provides basic computing resources
 - Operating system
 - Application programs
 - Users
- 

Four Components of a Computer System



Operating System Definition

- ▶ OS is a **resource allocator**
 - Manages all resources
 - Decides between conflicting requests for efficient and fair resource use
- ▶ OS is a **control program**
 - Controls execution of programs to prevent errors and improper use of the computer

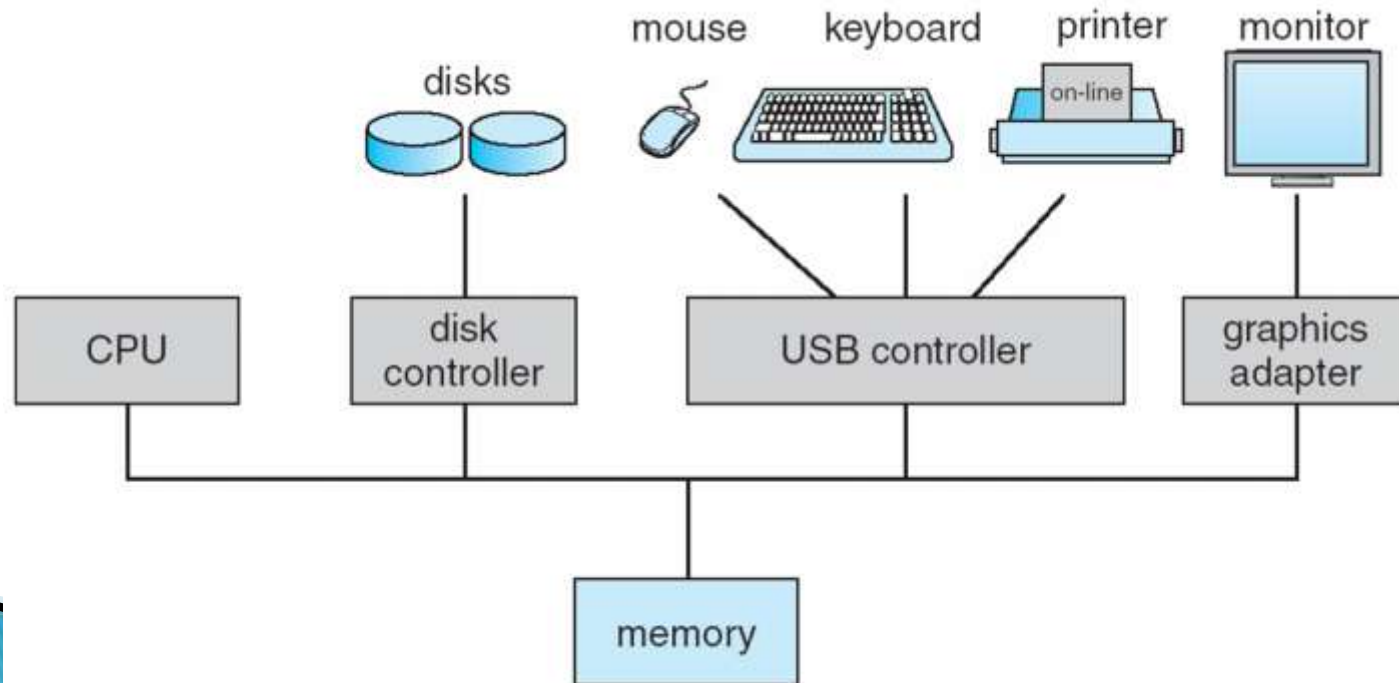
“The one program running at all times on the computer” is the **kernel**.

Computer Startup

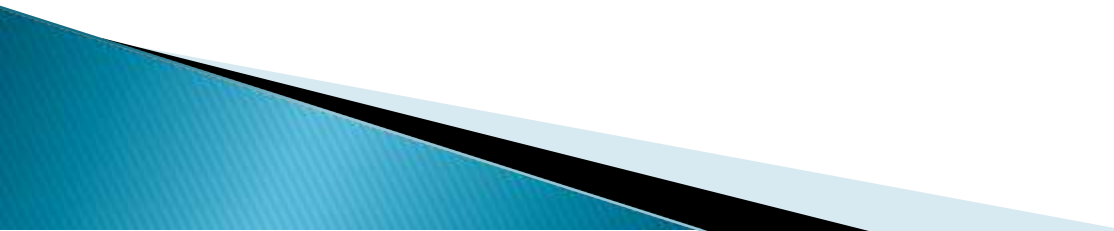
- ▶ **bootstrap program** is loaded at power-up or reboot
 - Typically stored in ROM or EPROM, generally known as **firmware**
 - Initializes all aspects of system
 - Loads operating system kernel and starts execution

Computer System Organization

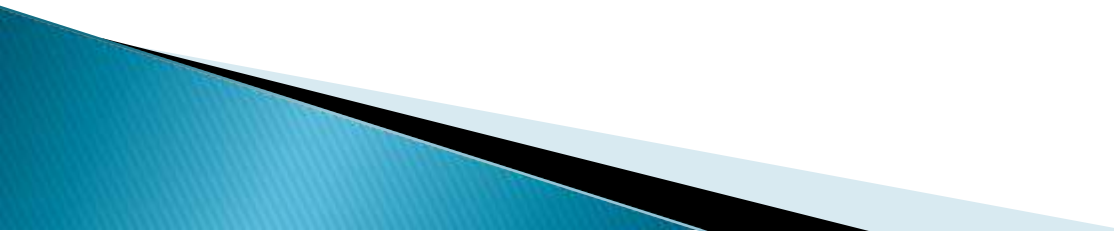
- ▶ Computer–system operation
 - One or more CPUs, device controllers connect through common bus providing access to shared memory
 - Concurrent execution of CPUs and devices competing for memory cycles



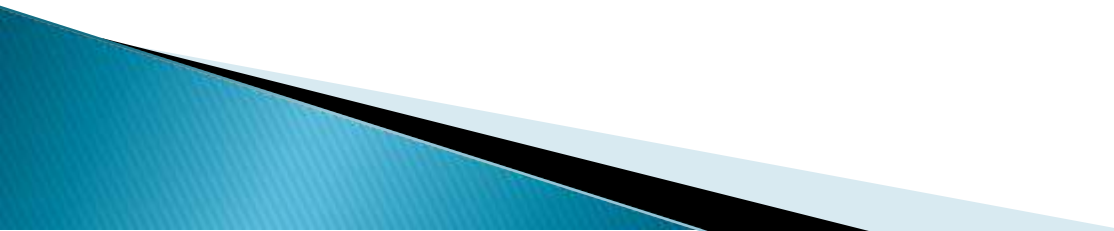
Computer–System Operation

- ▶ I/O devices and the CPU can execute concurrently
 - ▶ Each device controller is in charge of a particular device type
 - ▶ Each device controller has a local buffer
 - ▶ CPU moves data from/to main memory to/from local buffers
 - ▶ I/O is from the device to local buffer of controller
 - ▶ Device controller informs CPU that it has finished its operation by causing an **interrupt**
- 

Common Functions of Interrupts

- n Interrupt transfers control to the interrupt service routine generally, through the **interrupt vector**, which contains the addresses of all the service routines
 - n Interrupt architecture must save the address of the interrupted instruction
 - n A **trap** or **exception** is a software-generated interrupt caused either by an error or a user request
 - n An operating system is **interrupt driven**
- 

Interrupt Handling

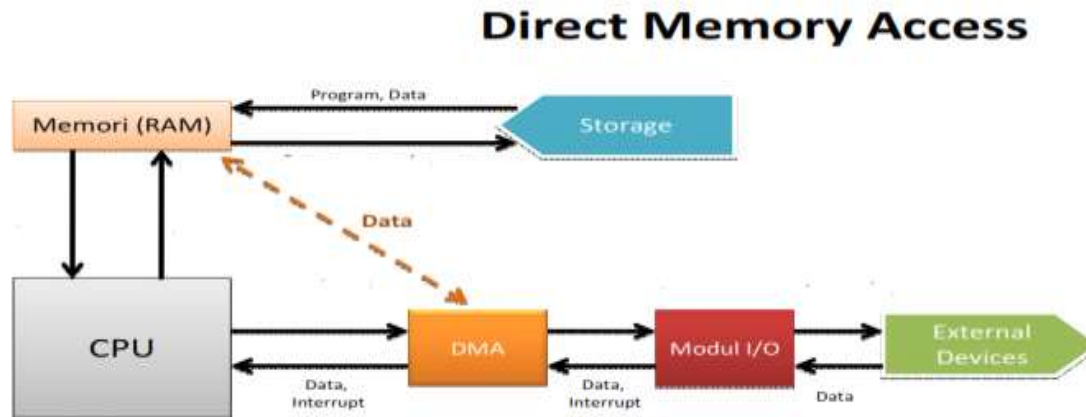
- ▶ The operating system preserves the state of the CPU by storing registers and the program counter
 - ▶ Determines which type of interrupt has occurred:
 - **polling**
 - **vectored** interrupt system
 - ▶ Separate segments of code determine what action should be taken for each type of interrupt
- 

I/O Structure

- ▶ After I/O starts, control returns to user program only upon I/O completion
 - Wait instruction idles the CPU until the next interrupt
 - Wait loop (contention for memory access)
 - At most one I/O request is outstanding at a time, no simultaneous I/O processing.
- ▶ After I/O starts, control returns to user program without waiting for I/O completion
 - **System call** – request to the OS to allow user to wait for I/O completion
 - **Device-status table** contains entry for each I/O device indicating its type, address, and state
 -

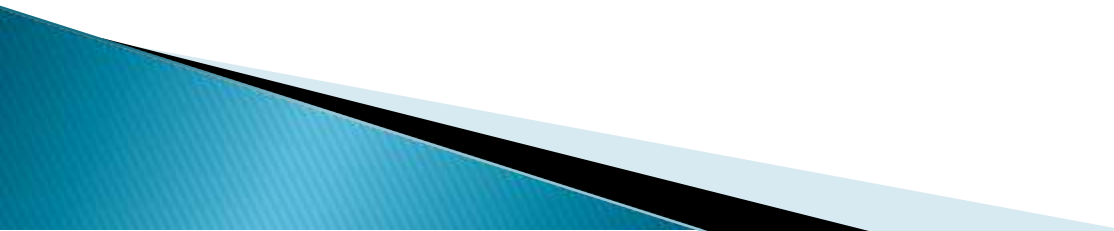
Direct Memory Access Structure

- ▶ Used for high-speed I/O devices able to transmit information at close to memory speeds



- ▶ Device controller transfers blocks of data from buffer storage directly to main memory without CPU intervention
- ▶ Only one interrupt is generated per block, rather than the one interrupt per byte


Storage Structure

- ▶ Main memory – only large storage media that the CPU can access directly
 - ▶ Secondary storage – extension of main memory that provides large **nonvolatile** storage capacity
 - ▶ Magnetic disks – rigid metal or glass platters covered with magnetic recording material
 - Disk surface is logically divided into **tracks**, which are subdivided into **sectors**
 - The **disk controller** determines the logical interaction between the device and the computer
- 

Storage Hierarchy

- ▶ Storage systems organized in hierarchy
 - Speed
 - Cost
 - Volatility
- ▶ **Caching** – copying information into faster storage system; main memory can be viewed as a cache for secondary storage
- ▶ **Device Driver** for each device controller to manage I/O
 - Provides uniform interface between controller and kernel

Storage-Device Hierarchy



Caching

- ▶ Important principle, performed at many levels in a computer (in hardware, operating system, software)
- ▶ Faster storage (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there




Computer–System Architecture

- ▶ **Multiprocessors** systems growing in use and importance
 - Also known as **parallel systems**, **tightly-coupled systems**
 - Advantages include:
 1. **Increased throughput**
 2. **Economy of scale**
 3. **Increased reliability – graceful degradation** or **fault tolerance**
 - Two types:
 1. **Asymmetric Multiprocessing**
 2. **Symmetric Multiprocessing**

Symmetric vs. Asymmetric Multiprocessing Architecture [1 / 2]


Symmetric Multiprocessing



(Shared Memory)

Vs


Asymmetric Multiprocessing



(No Shared Memory)


A Dual-Core Design

- ▶ **UMA** and **NUMA** architecture variations
- ▶ Multi-chip and **multicore**



Uniform memory access (UMA):
Each processor has uniform access to memory. Also known as **symmetric multiprocessors**, or SMPs (Sun E10000)


Non-uniform memory access (NUMA): Time for memory access depends on location of data. Local access is faster than non-local access. Easier to scale than SMPs (SGI Origin)



Clustered Systems

- ▶ Like multiprocessor systems, but multiple systems working together
 - Usually sharing storage via a **storage-area network (SAN)**
 - Provides a **high-availability** service which survives failures
 - **Asymmetric clustering** has one machine in hot-standby mode
 - **Symmetric clustering** has multiple nodes running applications, monitoring each other

Clustered Systems



Operating System Structure

- ▶ **Multiprogramming** needed for efficiency
 - ❖ Single user cannot keep CPU and I/O devices busy at all times
 - ❖ Multiprogramming organizes jobs (code and data) so CPU always has one to execute
 - ❖ A subset of total jobs in system is kept in memory
 - ❖ One job selected and run via **job scheduling**
 - ❖ When it has to wait (for I/O for example), OS switches to another job
- ▶ **Timesharing (multitasking)** is logical extension in which CPU switches jobs so frequently that users can interact with each job while it is running, creating **interactive** computing
 - ❖ **Response time** should be < 1 second
 - ❖ Each user has at least one program executing in memory \Rightarrow **process**
 - ❖ If several jobs ready to run at the same time \Rightarrow **CPU scheduling**
 - ❖ If processes don't fit in memory, **swapping** moves them in and out to run
 - ❖ **Virtual memory** allows execution of processes not completely in memory

Memory Layout for Multiprogrammed System

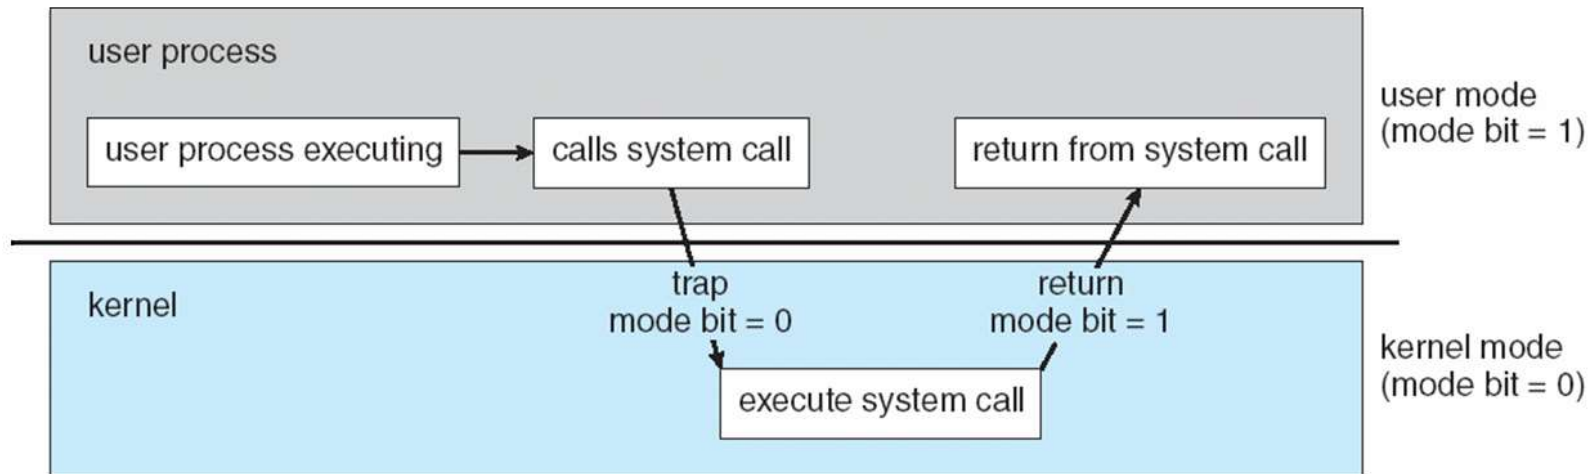


Operating-System Operations


- ▶ **Dual-mode** operation allows OS to protect itself and other system components
 - **User mode** and **kernel mode**
 - **Mode bit** provided by hardware
 - Provides ability to distinguish when system is running user code or kernel code
 - Some instructions designated as **privileged**, only executable in kernel mode
 - System call changes mode to kernel, return from call resets it to user
- ▶ Increasingly CPUs support multi-mode operations
 - i.e. **virtual machine manager (VMM)** mode for guest **VMs**

Transition from User to Kernel Mode

- ▶ Timer to prevent infinite loop / process hogging resources
 - Set interrupt after specific period
 - Operating system decrements counter
 - When counter zero generate an interrupt
 - Set up before scheduling process to regain

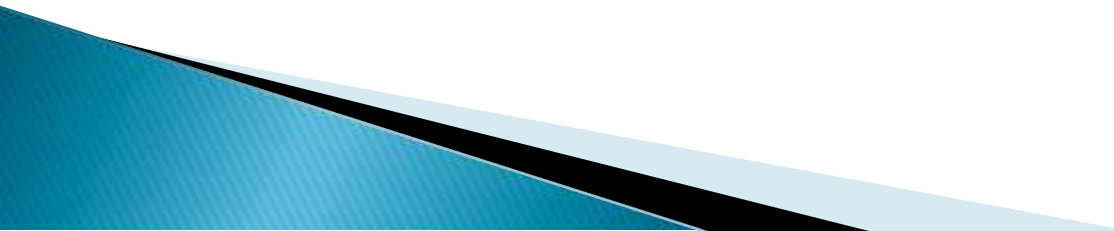


Process Management

- ▶ A process is a program in execution. It is a unit of work within the system. Program is a *passive entity*, process is an *active entity*.
 - ▶ Process needs resources to accomplish its task
 - CPU, memory, I/O, files
 - Initialization data
 - ▶ **Program counter (PC)**: Contains the address of an instruction to be fetched
 - ▶ Single-threaded process has one **program counter** specifying location of next instruction to execute
 - ▶ Multi-threaded process has one program counter per thread
- 

Process Management Activities

The operating system is responsible for the following activities in connection with process management:

- ▶ Creating and deleting both user and system processes
 - ▶ Suspending and resuming processes
 - ▶ process synchronization
 - ▶ process communication
 - ▶ deadlock handling
- 

Memory Management

- ▶ Memory management activities
 - Keeping track of which parts of memory are currently being used and by whom
 - Deciding which processes (or parts thereof) and data to move into and out of memory
 - Allocating and deallocating memory space as needed

Storage Management

▶ File-System management

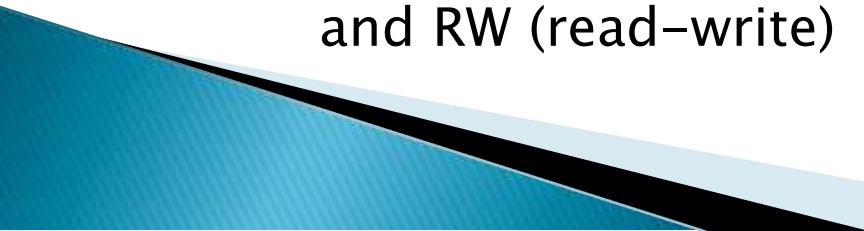
- Files usually organized into directories
- Access control on most systems to determine who can access what
- OS activities include
 - Creating and deleting files and directories
 - Primitives to manipulate files and dirs
 - Mapping files onto secondary storage
 - Backup files onto stable (non-volatile) storage media

Mass-Storage Management

▶ OS activities

- Free-space management
- Storage allocation
- Disk scheduling

▶ Some storage need not be fast

- Tertiary storage includes optical storage, magnetic tape
 - Still must be managed – by OS or applications
 - Varies between WORM (write-once, read-many-times) and RW (read-write)
- 

I/O Subsystem

- ▶ I/O subsystem responsible for
 - Memory management of I/O
 - caching (storing parts of data in faster storage for performance)
 - spooling (the overlapping of output of one job with input of other jobs)

Protection and Security

- ▶ **Protection** – any mechanism for controlling access of processes or users to resources defined by the OS
- ▶ **Security** – defense of the system against internal and external attacks
 - Huge range, including denial-of-service, worms, viruses, identity theft, theft of service


Computing Environments – Distributed

▶ Distributed

- Collection of separate, possibly heterogeneous, systems networked together
 - **Network** is a communications path,
 - **Local Area Network (LAN)**
 - **Wide Area Network (WAN)**
- **Network Operating System** provides features between systems across network


Computing Environments – Client-Server

- Client-Server Computing
 - Dumb terminals supplanted by smart PCs
 - Many systems now **servers**, responding to requests generated by **clients**

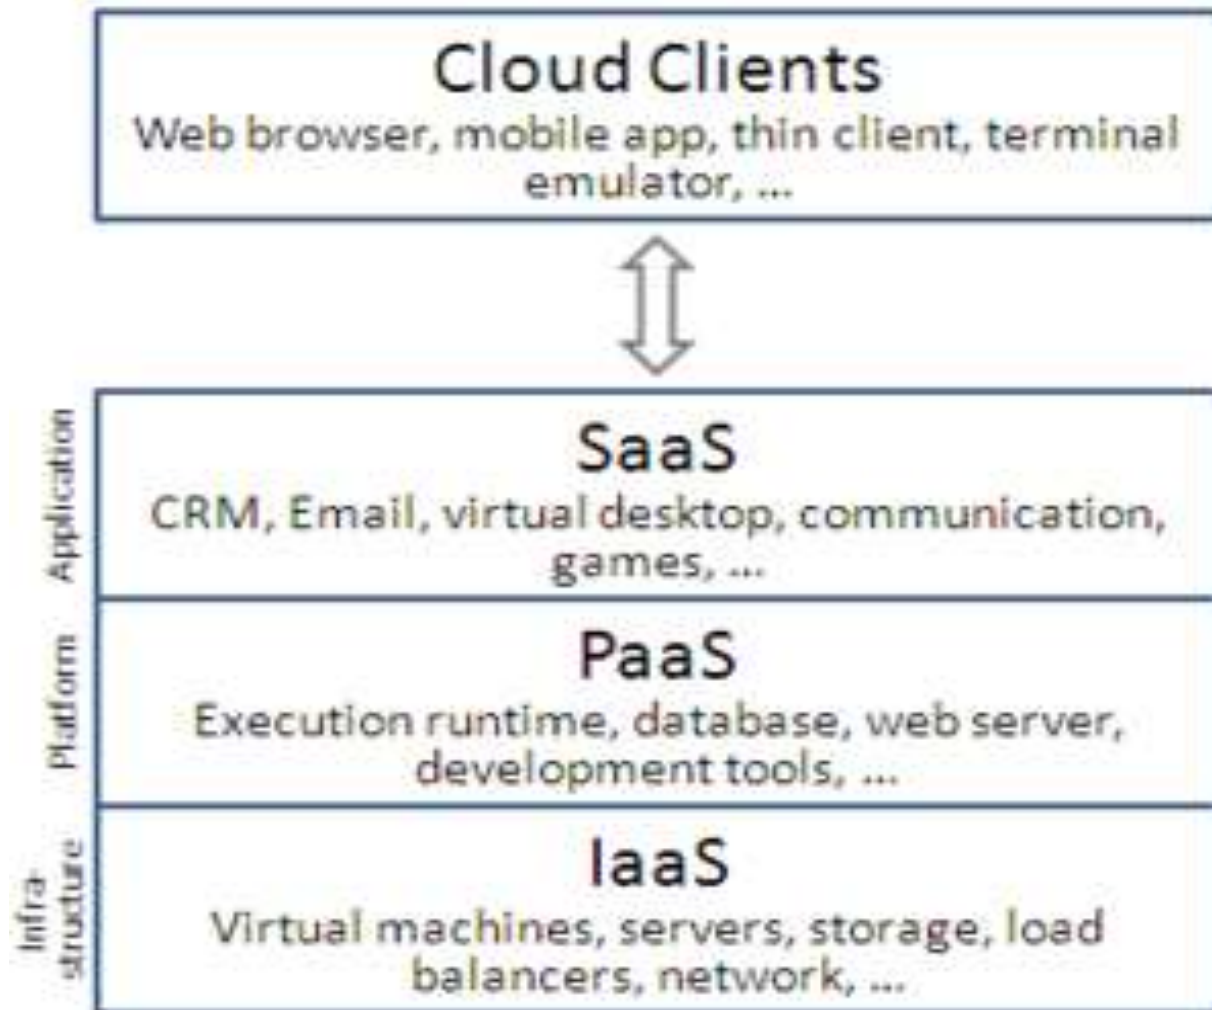


Computing Environments – Peer-to-Peer

- ▶ Another model of distributed system
- ▶ P2P does not distinguish clients and servers



Computing Environments – Cloud Computing



Computing Environments – Real-Time Embedded Systems

- ▶ Real-time embedded systems most prevalent form of computers
 - Vary considerable, special purpose, limited purpose OS, **real-time OS**
 - Use expanding
- ▶ Many other special computing environments as well
 - Some have OSes, some perform tasks without an OS
- ▶ Real-time OS has well-defined fixed time constraints
 - Processing *must* be done within constraint
 - Correct operation only if constraints met